

A REPORT
ON
RESPONSIBLE AI AND INTERPRETATION

Submitted by

Ms. Shaziya Banu

IS-2026-9761

INTERNSPARK

ARTIFICIAL INTELLIGENCE INTERNSHIP

Jul-2026

1. Introduction

Artificial Intelligence models are increasingly used in decision-making systems. While achieving high prediction accuracy is important, it is equally important to ensure that models are transparent, explainable, and fair.

Responsible AI focuses on understanding how machine learning models make predictions, identifying potential biases, and ensuring that predictions do not unfairly affect specific groups of users.

In this project, a Responsible AI analysis was conducted on the Customer Churn Prediction model developed in Task 1. Feature importance analysis, SHAP-based explainability, fairness assessment, and bias mitigation recommendations were performed.

2. Objective

The objectives of this project are:

- Understand the factors influencing customer churn predictions.
- Explain model predictions using SHAP (SHapley Additive exPlanations).
- Analyze fairness across demographic groups.
- Identify potential sources of bias.
- Propose practical mitigation strategies to improve responsible AI practices.

3. Dataset Description

The IBM Telco Customer Churn dataset was used for this analysis.

Dataset Characteristics

Attribute	Value
Total Records	7043
Total Features	21
Target Variable	Churn
Problem Type	Binary Classification

Target Variable

- Churn = Yes (Customer leaves the service)
- Churn = No (Customer remains with the service)

4. Data Preprocessing

The following preprocessing steps were performed before analysis:

4.1 Handling Missing Values

The TotalCharges column contained missing values.

Steps performed:

- Converted TotalCharges to numeric format.
- Replaced missing values using the median value.

4.2 Removing Irrelevant Features

The customerID column was removed because it does not contribute to customer churn prediction.

4.3 Encoding Categorical Variables


Categorical variables were converted into numerical values using Label Encoding to make them suitable for machine learning algorithms.

5. Model Used

A Logistic Regression classifier was used for the Responsible AI analysis.

Reasons for selecting Logistic Regression:

- High interpretability
- Strong performance in Task 1
- Easy integration with explainability techniques such as SHAP



The dataset was divided into:

- Training Data: 80%
- Testing Data: 20%


6. Feature Importance Analysis

Feature importance analysis was performed using the coefficients of the Logistic Regression model. The absolute coefficient values were used to determine which features have the strongest influence on churn prediction.

Observation

The analysis revealed that a small number of features contribute significantly to customer churn prediction, while others have relatively lower influence.

[7]:	Feature	Importance
5	PhoneService	1.020357
14	Contract	0.721993
15	PaperlessBilling	0.356723
8	OnlineSecurity	0.293026
11	TechSupport	0.256804
7	InternetService	0.227226
3	Dependents	0.184115
1	SeniorCitizen	0.178126
9	OnlineBackup	0.153567
10	DeviceProtection	0.082099



7. Model Explainability Using SHAP

SHAP (SHapley Additive exPlanations) was used to understand how individual features contribute to model predictions.

SHAP provides:

- Global Interpretability
- Local Interpretability
- Feature Contribution Analysis


7.1 SHAP Summary Plot

The SHAP summary plot visualizes how each feature affects model predictions across all customers.

Features with higher SHAP values have a stronger influence on the prediction outcome.

Observation

The plot clearly identifies the most influential features driving churn predictions and shows whether feature values increase or decrease churn probability.




7.2 SHAP Feature Importance Plot

A SHAP bar plot was generated to rank features according to their overall contribution.

Observation

The SHAP ranking confirms the importance of several customer-related attributes that strongly influence churn behavior.



8. Fairness and Bias Analysis

Responsible AI requires evaluating whether model behavior differs across demographic groups.

Two demographic attributes were analyzed:

1. Gender
2. Senior Citizen Status


8.1 Gender-Based Fairness Analysis

Customer churn rates were compared across gender groups.

Observation

The churn rates for male and female customers were examined to identify any significant disparities.

A large difference in churn rates may indicate the need for further fairness evaluation.




8.2 Senior Citizen Fairness Analysis

Customer churn rates were analyzed for:

- Senior Citizens
- Non-Senior Citizens

Observation

The comparison helps identify whether the model behaves differently for customers of different age groups.



9. Bias Assessment

Based on the fairness analysis:

- No direct evidence of intentional bias was identified.
- Differences observed between demographic groups may arise from genuine variations within the dataset.
- Continuous monitoring is required to ensure fairness during deployment.

Bias evaluation should be considered an ongoing process rather than a one-time activity.

10. Bias Mitigation Recommendations

The following recommendations can help reduce the risk of unfair outcomes:

1. Balanced Training Data

Ensure that all demographic groups are adequately represented during model training.

2. Continuous Fairness Monitoring

Regularly evaluate model performance across demographic groups after deployment.

3. Explainable AI Integration

Use explainability techniques such as SHAP to understand model behavior and identify unexpected patterns.

4. Periodic Model Retraining

Retrain the model using updated and representative datasets to prevent performance degradation and bias accumulation.

5. Fairness Evaluation Before Deployment

Incorporate fairness testing as a mandatory step before production deployment.

11. Results Summary

The Responsible AI analysis successfully provided:

- Feature Importance Analysis
- SHAP-Based Explainability
- Gender Fairness Evaluation
- Senior Citizen Fairness Evaluation
- Bias Assessment
- Practical Mitigation Recommendations

These analyses improve transparency and trustworthiness of the customer churn prediction model.

12. Conclusion

In this project, Responsible AI techniques were applied to a Customer Churn Prediction model.

Feature importance analysis and SHAP explainability helped understand the factors influencing churn predictions. Fairness analysis was conducted across gender and senior citizen groups to identify potential bias.

The study demonstrated that explainability and fairness evaluation are essential components of machine learning systems. Incorporating these practices improves transparency, accountability, and trust in AI-driven decision-making.

Technologies Used

- Python 3.11
- Pandas
- NumPy

- Matplotlib
- Seaborn
- Scikit-Learn
- SHAP
- Jupyter Notebook

References

1. SHAP Documentation
2. Scikit-Learn Documentation
3. IBM Telco Customer Churn Dataset
4. Pandas Documentation
5. NumPy Documentation